

Implement a small search engine. The main steps are:

- Read in the tokenized and stemmed document collection provided in the file `tccorpus.txt` inside [tccorpus.zip](#). This is an early standard collection of abstracts from the *Communications of the ACM*.
- Build a simple inverted indexer that reads the corpus and writes the index. You should invoke it with

```
indexer tccorpus.txt index.out
```

- Implement the BM25 ranking algorithm and write a program to provide a ranked list of documents for a file with one or more queries. You should pass parameters for the index file, the query file, and the maximum number of document results, and return documents on the standard output, like so:

```
bm25 index.out queries.txt 100 > results.eval
```

- Submit the output from this run, with the top 100 document IDs and their BM25 scores for each test query according to the following format:

```
query_id Q0 doc_id rank BM25_score system_name
```

The string `Q0` is a literal used by the evaluation script. You can use any space-free token for your `system_name`.

- Also, submit your code, instructions for compiling it, and a short report describing your implementation.

Tokenized Document Collection

- The provided `tccorpus.txt` file is in the format:
 - A `#` followed by a document ID
 - Lines below the document ID line contain stemmed words from the document.
- For example:

```
# 1
this is a tokenized line for document 1
this is also a line of document 1
# 2
from here lines for document 2 begin
...
...
# 3
...
```

Building an Inverted Index: The following data structures are required for BM25 computation:

- Term frequencies (tf) are stored in inverted lists: `word -> (docid, tf), (docid, tf), ...`
- For this assignment, you don't need to consider term positions within documents.
- Store the number of tokens in each document in a separate data structure.
- You may employ any concrete data structures convenient for the programming language you are using, as long as you can write them to disk and read them back in when you want to run some queries.

BM25 Ranking

1. Retrieve all inverted lists corresponding to terms in a query.
2. Compute BM25 scores for documents in the lists.
3. Make a score list for documents in the inverted lists.
4. Accumulate scores for each term in a query on the score list.
5. Assume that no relevance information is available.
6. For parameters, use $k_1=1.2$, $b=0.75$, $k_2=100$.
7. Sort the documents by the BM25 scores.

Test Queries: Use the following stemmed test queries, also provided in the file [queries.txt](#):

Query ID	Query Text
1	portabl oper system
2	code optim for space effici
3	parallel algorithm
4	distribut comput structur and algorithm
5	appli stochast process
6	perform evalu and model of comput system
7	parallel processor in inform retriev